



Programming Language Translation

Practical 2: week beginning 23 July 2018

Hand in this prac sheet before lunch time on your next practical day, correctly packaged in a transparent folder with your solutions and the "cover sheet". Unpackaged and late submissions will not be accepted -- you have been warned. Since the practical will have been done on a group basis, please hand in one copy of the cover sheet for each member of the group. These will be returned to you in due course, signed by the marker.

Objectives

In this practical you are to

- become familiar you with the workings of two simple machine emulators for the PVM pseudo-machine that we shall use frequently in the course.
- gain some experience with the machines, writing machine code for them, comparing them and extending them.

Copies of this handout, the cover sheet and the Parva language report are available on the RUConnected course page.

Outcomes

When you have completed this practical you should understand:

- the opcode set for the Parva Virtual Machine (PVM);
 - how to write and debug machine level code for the PVM;
 - how to extend the PVM to incorporate new opcodes;
 - why, and by how much, interpretive systems can vary in execution overhead.
-

To hand in (45 marks)

This week you are required to hand in, besides the cover sheet:

- Listings of the final version of the assembler/emulator system you produce or (preferably) extracts showing only the extensions clearly (to save paper!), and your solutions to the programming exercises below. Use LPRINT, please. *One listing/group please.*
- Additionally, electronic copies of source code for those exercises, using the electronic submission system.
- Discussion of the experiments in Tasks 3, 4 and 9.

Keep the cover sheet and your solutions until the end of the semester. Check carefully that your mark has been entered into the Departmental Records.

You are referred to the rules for practical submission which are clearly stated in our Departmental Handbook. However, for this course pracs must be posted in the "hand-in" box inside the laboratory before the next practical session and not given to demonstrators during the session.

A rule not stated there, but which should be obvious, is that you are not allowed to hand in another student's or group's work as your own. Attempts to do this will result in (at best) a mark of zero and (at worst) severe disciplinary action and the loss of your DP. You are allowed -- even encouraged -- to work and study with other students, but if you do this you are asked to acknowledge that you have done so on *all* cover sheets and with suitable comments typed into *all* listings. You are expected to be familiar with the University Policy on Plagiarism.

Task 1 Creating a working directory and unpacking the prac kit

There are several files that you need, zipped up this week in the file `PRAC2.ZIP`.

Immediately after logging on, get to the DOS command line level and log onto your file space.

Copy the zipped prac kit into a new directory/folder in your file space, either from the server (I:) or RUConnected.

```
md  prac2
cd  prac2
copy i:\csc301\trans\prac2.zip
unzip prac2.zip
```

Task 2 Build the assemblers

In the working directory you will find C# files that give you two minimal assemblers and emulators for the PVM stack machine (described in Chapter 4.7). These files have the names

<code>PVMAsm.cs</code>	a simple assembler
<code>PVMPushPop.cs</code>	an interpreter/emulator, making use of auxiliary Push and Pop methods
<code>PVMInLine.cs</code>	an interpreter/emulator, with the pushing and popping "inlined"
<code>Assem.cs</code>	a driver program

`PVMPushPop` incorporates rather more constraint checking than is found in `PVMInLine`, and also has an option for doing a line-by-line trace of the code it is interpreting.

You compile and make two nominally equivalent assembler/interpreter systems by issuing the batch commands

```
MAKEASM1  make up a system ASM1.EXE using PVMPushPop as the PVM
MAKEASM2  make up a system ASM2.EXE using PVMInLine as the PVM
```

These take as input a "code file" in the format shown in the examples in Section 4.5 and in the prac kit. Make up the minimal assembler/interpreters and, as a start, run these using a supplied small program:

ASM1 lsmall.pvm this will prompt you for input and output files and for tracing options
 ASM2 lsmall.pvm
 ASM2 lsmall.pvm immediate this will enter the emulator immediately after compiling

Wow! Isn't Science wonderful? Try the interpretation with and without the trace option, and familiarize yourself with the trace output and how it helps you understand the action of the virtual machine (ASM1 only).

Task 3 A look at PVM code (1+2+2 = 5 marks)

Consider the following gem of a Parva program which reads a list of integers and writes it in reverse.

```
void Main () {
// Read a zero-terminated list of numbers and write it backwards
// P.D. Terry, Rhodes University, 2015
const max = 10;
int[] list = new int[max];
int i = 0, n;
read(n);
while ((n != 0) && (i < max)) { // input loop
    list[i] = n;
    i++;
    read(n);
}
while (i > 0) { // output loop
    i--;
    write(list[i]);
}
} // Main
```

You can compile and run this (PARVA REVERSE.PAV) at your leisure to make quite sure that it works.

The Parva compiler supplied to you this week is not the same as last week -- it only allows a single `Main()` function, but it includes "else" and the modulo "%" operator, supports a "repeat" ... "until" statement, and allows increment and decrement operations like `i++` and `array[j]--` (not within expressions, however).

In the prac kit you will also find a translation of this program into PVM code (REVERSE.PVM). Study this code and complete the following tasks:

0 DSP	3	42 LDA	1
2 LDA	0	44 LDA	1
4 LDC	10	46 LDV	
6 ANEW		47 LDC	1
7 STO		49 ADD	
8 LDA	1	50 STO	
10 LDC	0	51 LDA	2
12 STO		53 INPI	
13 LDA	2	54 BRN	16

15	INPI		56	LDA	1
16	LDA	2	58	LDV	
18	LDV		59	LDC	0
19	LDC	0	61	CGT	
21	CNE		62	BZE	84
22	LDA	1	64	LDA	1
24	LDV		66	LDA	1
25	LDC	10	68	LDV	
27	CLT		69	LDC	1
28	AND		71	SUB	
29	BZE	56	72	STO	
31	LDA	0	73	LDA	0
33	LDV		75	LDV	
34	LDA	1	76	LDA	1
36	LDV		78	LDV	
37	LDXA		79	LDXA	
38	LDA	2	80	LDV	
40	LDV		81	PRNI	
41	STO		82	BRN	56
			84	HALT	

- (a) How can you tell that the translation has not used short-circuit Boolean operations?
- (b) Add commentary to the code that "matches" the Parva code fairly closely. Have a look at the `LSMALL.PVM` code example in the prac kit to see a "preferred" style of commentary, where the high level code appears as commentary on the low level code.
- (c) What would you need to change if you wanted to make use of short-circuit Boolean operations? (You should test your ideas with the first of the two assemblers, ASM1).

The (modified and suitably commented) REVERSE.PVM file must be submitted for assessment as well as the answer to question(a).

Task 4 Execution overheads - part one (3 marks)

In the prac kit you will find a translation `SIEVE1.PVM` of a cut down version of a prime-counting program `SIEVE.PAV` based on last week's exercises (the source code is also there, but is not printed here to save paper).

Run `SIEVE1.PVM` through both versions of the assemblers and obtain timings for a suitable upper limit (say 4000) and number of iterations (say 100) for the combinations:

Hint: The lab computers are very fast. You may have to alter those suggestions quite a bit to produce measurably distinct timings. Do this simply by changing the appropriate constants in the PVM file before you assemble it

Comment on the results. Are they what you expect? If not, why not?

Submit your discussion as a short essay.

Task 5 Coding the hard way (6 marks)

Time to do some creative work at last. Task 5 is to produce an equivalent program to the Parva one below (PALIN.PAV), but written directly in the PVM stack-machine language (PALIN.PVM). In other words, "hand compile" the Parva algorithm directly into the PVM machine language. You may find this a bit of a challenge, but it really is not too hard, just a little tedious, perhaps.

```
void Main () {
// Read a sequence of numbers and report whether they form a palindromic
// sequence (one that reads the same from either end)
// Examples:  1 2 3 4 3 2 1  is palindromic
//           1 2 3 4 4 3 2  is non-palindromic
// P.D. Terry, Rhodes University, 2015

int
    n,                      // number of items
    low, high,              // indices of items to be compared
    item;                   // latest item read
bool
    isPalindrome;           // Boolean flag
int [] list = new int [10]; // the list of items

n = 0;
read(item);
while (item != 0) {
    list[n] = item;
    n = n + 1;
    read(item);
}
isPalindrome = true;        // optimist
low = 0; high = n - 1;      // initial indices
while (low < n - 1) {       // sweep through the list
    if (list[low] != list[high])
        isPalindrome = false; // bad luck
    low = low + 1; high = high - 1; // adjust indices
}
if (isPalindrome) write("Palindromic sequence");
else write("Non-palindromic sequence");
} // Main
```

Health warning: if you get the logic of your program badly wrong, it may load happily, but then go berserk when you try to interpret it. You may discover that the interpreter is not so "user friendly" as all the encouraging remarks in the book might have led you to believe interpreters all to be. Later we may improve it quite a bit. (Of course, if your machine-code programs are correct you won't need to do so. As has often been said: "Any fool can write a translator for source programs that are 100% correct".)

The most tedious part of coding directly in PVM code is computing the destination addresses of the various branch instructions.

Hint: As a side effect of assembly, the ASM system writes a new file with a .COD extension showing what has been assembled and where in memory it has been stored. Study of a .COD listing will often give you a good idea of what the targets of branch instructions should really be.

The (suitably commented) PALIN.PVM file must be submitted for assessment.

Task 6 Trapping overflow and other pitfalls (6 marks)

Several of the remaining tasks in this prac require you to examine the machine emulator to learn how it really works, and to extend it to improve some opcodes and to add others.

In the prac kit you will discover two programs deliberately designed to cause chaos. `DIVZERO.PVM` bravely tries to divide by zero, and `MULTBIG.PVM` embarks on a continued multiplication that soon goes out of range. Try assembling and interpreting them with both systems to watch disaster happen.

Now we can surely do better than that! Modify the interpreters (`PVMPushPop.cs` and `PVMInLine.cs`) so that they will anticipate division by zero or multiplicative overflow, and change the program status accordingly, so that users will be told the errors of their ways and not left wondering what has happened.

You will have to be subtle about this -- you have to detect that problems are going to occur *before* things "go wrong", and you must be able to detect it for negative as well as positive overflow conditions.

The suggested program in Task 5 has a very small working array. What happens if you try to execute `PALIN.PVM` with a list of, say, 12 numbers? Try this with both assemblers -- and then fix the broken one!

Hint: After you edit any of the source code for the assemblers you will have to issue the `MAKEASMx` commands to recompile them, of course. It's easy to forget to do this and then wonder why nothing seems to have changed.

Task 7 Support for characters (10 marks)

If the PVM and Parva could only handle characters as well as integers and Booleans, we could write a program like the exciting one below that reads a string of characters terminated with a period (full stop) and then encrypts it in lower case, using the "rot 13" algorithm. (`ENCOD.PAV`).

```
void Main() {
// rot13 encryption of a text terminated with a period
// P.D. Terry, Rhodes University, 2015

char ch;
repeat {
    read(ch);
    ch = lower(ch);
    if (isLetter(ch)) ch = (char) ('a' + (ch - 'a' + 13) % 26);
    write(ch);
}
until (ch == '.');
```

```
} // Main
```

Not a problem for the assembler system. All we need to do is add appropriate opcodes to our virtual machine -for a start, `INPC` for reading a character and `PRNC` for writing a character - to open up exciting possibilities.

Hint: Adding "instructions" to the pseudo-machine is easy enough, but you must be careful to make sure you modify all the parts of the system that need to be modified. Before you begin, study the code in the definition of the stack machine carefully to see where and how the opcodes are defined, how they are mapped to the mnemonics, and in which switch/case statements they are used.

This example has implied the availability of a method for converting characters to lowercase, which is easily added to the PVM by introducing a special opcode. We have also hinted at the desirability of supporting the infamous `++` and `--` operators, which can be handled by special opcodes that take less space (and should take less time to execute) than the tedious sequences needed for code corresponding directly to an assignment statement like `n = n + 1;` as seen before.

Extend the machine and the assembler still further with opcodes `ISLET`, `LOW`, `INC` and `DEC`, and hand compile the `rot13` program to use them.

Submit your hand-coded `rot13.pvm`.

Hint: Be careful. Think ahead! Don't limit your `INC` and `DEC` opcodes to cases where they can handle assignment statements like `x++`; only. In some programs you might want to have assignment statements like `List[N+6]++`; . Regard these as statements, and not as components of expressions, as they are in C#.

How do you decode a message that has been encrypted by this program?

Task 8 Improving the opcode set still further (11 marks)

Section 4.10 of the text discusses the improvements that can be made to the system by adding new single-word opcodes like `LDC_0` and `LDA_0` in place of double-word opcodes for frequently encountered operations like `LDC 0` and `LDA 0`, and for using load and store opcodes like `LDL N` and `STL N` (and, equivalently, opcodes like `LDL_0` and `STL_0` for frequently encountered special cases).

Enhance both versions of your PVMs to incorporate the following opcodes:

<code>LDL N</code>	<code>STL N</code>		
<code>LDA_0</code>	<code>LDA_1</code>	<code>LDA_2</code>	<code>LDA_3</code>
<code>LDL_0</code>	<code>LDL_1</code>	<code>LDL_2</code>	<code>LDL_3</code>
<code>STL_0</code>	<code>STL_1</code>	<code>STL_2</code>	<code>STL_3</code>
<code>LDC_0</code>	<code>LDC_1</code>	<code>LDC_2</code>	<code>LDC_3</code>

Hint: Several of the above are very similar to one another. Note that the assemblers have already been primed with the mappings from these mnemonics to integers, but, once again, you must be careful to make sure you modify all the parts of the system that need extending - you will have to add quite a bit to various switch statements to complete the tasks. Do this for both versions of the PVM.

Try out your systems by developing an "improved" version of `PALIN.PVM`, say `PALINC.PVM` that handles sentences, not numbers, on the lines of

```
void Main () {
// Read a sequence of characters terminated by a period and report whether
// they form a palindrome (one that reads the same from either end)

// Examples:   too hot to hoot.   is palindromic
//             1234432.           is non-palindromic
// P.D. Terry, Rhodes University, 2015

int
    n,                      // number of characters
    low, high;              // indices of characters to be compared
char
    ch;                     // latest character read
bool
    isPalindrome;           // Boolean flag
char [] str = new char [100]; // the string to be checked

n = 0;
read(ch);
while (ch != '.') {
    if (ch > ' ') {          // effectively ignore spaces etc
        str[n] = lower(ch); // convert to lower case
        n++;
    }
    read(ch);
}
isPalindrome = true;        // optimist
low = 0; high = n - 1;      // initial indices
while (low < n - 1) {       // sweep through the string
    if (str[low] != str[high])
        isPalindrome = false; // bad luck
    low++; high--;          // adjust indices
}
if (isPalindrome) write("Palindromic string");
else write("Non-palindromic string");
} // Main
```

The final assembler/emulators must be submitted for assessment, as must `PALINC.PVM`. It would help if you simply printed only those parts of the interpreters that you have modified in tasks 6, 7 and 8 -- large portions of the original will not need to change at all. Be careful to include the sections that deal with the run-time error trapping, however.

Task 9 Execution overheads - part two (4 marks)

You might think it is pretty obvious that using as many `STL` and `LDL` opcodes as possible should make your programs smaller, faster, better. Experiment to see whether this is true and, if so, how big this effect is.

In the prac kit you will find further translations `SIEVE2.PVM` and `SIEVE3.PVM` of the same prime-counting program `SIEVE.PAV` as was used in Task 4, but this time using the extended opcode set developed in the last task in various ways.

Run `SIEVE2.PVM` and `SIEVE3.PVM` through both versions of your modified assemblers and obtain timings for the same limit (say 4000) and number of iterations (say 100) as in Task 4.

Hint: You may have to alter those suggestions quite a bit to produce measurably distinct timings.

Comment on the results. Are they what you expect? If not, can you suggest why not?

Hopefully by now you will have found that interpreters are quite easy to develop, but this prac should show you that they are not necessarily very "efficient". What changes could one make to improve the efficiency of the interpreter for the PVM still further? (If you are very keen you might try out some of your ideas, but I suppose that is wishful thinking. Sigh ...)

Think carefully about all this. Have fun, and good luck.